Real-time Rendering of 3D "Fractal-like" Geometry

Deliverable 1: Final Year Dissertation

22/11/2021

by

Solomon Baarda

Meng Software Engineering

Heriot-Watt University

Supervisor: Dr Benjamin Kenwright

Second Reader: Ali Muzaffar

Abstract

A fractal is a recursively created never-ending pattern that is usually self-similar. Separate from Euclidean geometry, fractal geometry describes the more non-uniform shapes found in nature, like clouds, mountains, and coastlines. Fractal patterns exist everywhere in the universe, whether we can see them or not. From DNA molecules to the structure of galaxies, and everything in between. Fractals patterns appear everywhere in nature and many technological breakthroughs have been made through studying them.

With the increasing popularity in fractal geometry, and increasing computing power, fractal rendering software has become far more common in the last decade. However, only a small number of these programs are capable of rendering 3D fractals in real time, and those that are capable, are mostly written using graphics shaders which contain lots of code duplication between scenes. This makes it hard for a beginner to get into rendering 3D fractals as they must be competent in the chosen shader language and must understand the complex theory of rendering fractals. This project aims to produce a real-time 3D fractal geometry renderer, for which it is easy for a user to create new scenes and add geometry to it.

Declaration

I, Solomon Baarda confirm that this work submitted for assessment is my own and is expressed in

my own words. Any uses made within it of the works of other authors in any form (e.g., ideas,

equations, figures, text, tables, programs) are properly acknowledged at any point of their use. A list

of the references employed is included.

Signed: *Solomon Baarda*

Date: 16/11/2021

Table of Contents

Table of Figures

Table of Tables

Common Definitions

Table 1.1.i Common definitions

Word	Definition
Complex number	Number system containing an imaginary unit, defined as $i = \sqrt{-1}$
Convex polyhedron	3D equivalent of a regular 2D polygon
Euclidian geometry	A geometry containing basic objects like points and lines
Fractal	Recursively created never-ending pattern that is usually self-similar
Frame	One of many still images that make up a moving image
Geometry	Branch of mathematics concerned with the properties of space, distance, shape, size, and positions
Method overriding	Object-oriented programming technique which allows a subclass to change the implementation of a method that a parent class is providing
Object-oriented programming	Programming style that organises its software design around reusable objects
Polygon	2D shape with three or more sides
Polyhedrons	3D shape with six or more faces
Quaternion	Extension of the complex numbers, often used for storing position, translation, and scale values in 3D space
Ray	Line in 3D space with a beginning and an end
Regular polygon	Shape where all edges are the same length, and all angles between vertices are also equal
Render	Process of creating a still image using a computer
Vector	Mathematical object representing a position in space relative to another

Common Abbreviations

Table 1.1.ii Common abbreviations

Word	Abbreviation
CPU	Central Processing Unit
DE	Distance estimation function
FPS	Frames per second
GPU	Graphics Processing Unit
PC	Personal computer
SDF	Signed distance function

1 INTRODUCTION

1.1 PROJECT DESCRIPTION

A fractal is a recursively created never-ending pattern that is usually self-similar [1]. Separate from Euclidean geometry, fractal geometry describes the more non-uniform shapes found in nature, like clouds, mountains, and coastlines. Beniot Mandelbrot, inventor of the concept of fractal geometry, famously wrote "Clouds are not spheres, mountains are not cones, coastlines are not circles, and bark is not smooth, nor does lightning travel in a straight line" [2]. Fractal patterns exist everywhere in our lives [3], whether we can see them or not. From DNA molecules to the structure of galaxies, and everything in between.

Fractals have many applications in the real world [1]–[4]. In medicine, fractals have been used to help distinguish between cancerous cells which grow abnormally, and healthy human blood vessels which typically grow in fractal patterns. In fluid mechanics, fractals have been used to help model both complex turbulence flows and the structure of porous materials. In computer science, fractal compression is an efficient method for compressing images and other files and uses the fractal characteristic that parts of a file will resemble other parts of the same file. Fractal patterns are also used in the design of some cell phone and Wi-Fi antennas, as the fractal design allows them to be made more powerful and compact than other designs. Even losses and gains in the stock market have been described in terms of fractal mathematics. All these applications of fractals demonstrate that this geometry is a fundamental part of our universe, both in physical objects and in theoretical concepts.

With the increasing popularity in fractal geometry, and increasing computing power, fractal rendering software has become far more common in the last decade [5]. However, only a small number of these programs are capable of rendering 3D fractals in real time, and those that are capable, are mostly written using graphics shaders which contain lots of code duplication between scenes. This makes it hard for a beginner to get into rendering 3D fractals as they must be competent in the chosen shader

language and understand the complex theory of rendering fractals. This purpose of this project, therefore, is to design and implement a real-time 3D fractal geometry renderer, for which it is easy for a user to create and add geometry to scenes, in the hopes that it will fill the gap in the market for entry-level 3D fractal viewing software.

The term fractal has been used throughout this report to describe a pattern or geometry that displays the recursive self-similarity characteristic of fractals. The term fractal-like has been used to describe something that appears to display the fractal characteristics but may not truly contain infinite detail.

1.2 AIMS & OBJECTIVES

The aim of this project is to develop an application which can render 3D fractal geometry in real-time. The render will include common optical effects such as ambient occlusion, soft shadows, and lighting. In addition, it should be possible to create scenes and view them using the application, and this process should be as straight forward as possible.

Listed below are the key objectives that this project sets out to achieve. These objectives will help guide the project in the correct direction to achieve its aim.

Objective 1: Research topic

Background research of the project area to gain a better understanding of the chosen topic and the scope of the project. Once this first stage has been completed, the research must be kept up to date and any significant developments in the project research area should be explored.

Objective 2: Investigate existing solutions

Several relevant existing solutions exist, and an analysis of their strengths and flaws would benefit the project. This information will be kept up to date and any relevant newly released applications will be added and reviewed.

Objective 3: Core functionality

Implement the core functionality of a non-real-time 3D fractal renderer. This objective forms the safe core of the project, and the following objectives build upon this.

Objective 4: Additional functionality

This involves adding additional functionality to the renderer, such as making the application capable of real-time rendering, adding a game-loop, adding a controllable camera, making scenes dynamic, and adding optical effects and lighting. The scope of this objective can be increased or decreased as necessary, and several additional stretch goals have been included in the requirements specification.

Objective 5: Evaluation

Benchmark the performance of the application across various systems and evaluate how successfully the project aim was achieved.

Objective 6: Create user documentation

Create documentation to assist users or future developers of the application.

These objectives form the main tasks to be completed during the duration of this project, and the requirements specification in section 3.1 and the Gantt chart in section 6.5 have been structured around these. It is necessary to complete some of the objectives in the order they are specified, as they build upon previous objectives. Objectives one and two will run for the entire duration of the project, to ensure that the project stays up to date with current developments. Objectives three to six, however, relate to specific functionality to be implemented, and must be completed in order.

These objectives have been created bearing the SMART properties in mind. SMART stands for Specific, Measurable, Achievable, Realistic and Time constrained.

1.3 SCOPE

The scope of the project has been carefully considered, and several stretch goals have been included in the requirements specification if good progress is made. Objective 4 leaves large amounts of flexibility and can be extended or cut back depending on time constraints.

At the time of writing this report, the first stage of objectives one and two have already been completed. Additionally, progress has been made towards objective three and initial experimentation rendering still images of the Mandel bulb fractal and other geometry has been successful. Some of

these renders can be viewed in the appendix, section 9.1, and all project files can be found in the project GitHub repository [6]. In addition, some experimentation with OpenCL, a library that allows code to be executed in parallel on the graphics card has been completed. These experiments were done to gain familiarity with this style of programming in the hope to reduce the learning curve of this new software.

1.4 DOCUMENT STRUCTURE

Continuing from section 1, section 2 discusses relevant background literature for the project. Section 3 contains the requirements specification for the application and Section 4 discusses the technical design of the application. Section 5 discusses the strategy for testing and evaluating the application. Section 6 contains a plan for the project and Section 7 concludes the document.

2 LITERATURE REVIEW

This literature review contains a breakdown of information relevant for understanding the complexity of this project and contains details of how the problem will be tackled. While this review contains explanations of all relevant key concepts, basic knowledge of recursion, vector maths, and complex numbers is assumed.

This review is split into several sections, first the theory of 2D fractals and their 3D counterparts will be discussed. Then the key concepts of ray marching, the chosen method of rendering fractals, will be outlined. This will be followed by a brief introduction to the core concepts of GPU parallel programming, which is followed by a short analysis of several relevant existing solutions.

As discussed in the introduction, a fractal is a recursively created never-ending pattern that is usually self-similar [1]. This concept defines fractal geometry, which describes up the more non-uniform shapes found in nature, like clouds, mountains, and coastlines. There exist several ways of artificially creating a fractal, from manually defining a simple repeating pattern to studying the convergence of

equations. This section will discuss several common 3D fractals (and their 2D counterparts) and the methods used for creating them.

2.1.1 Sierpiński Tetrahedron

The Sierpiński tetrahedron, also known as the Sierpiński pyramid, is a 3D representation of the famous 2D Sierpiński tringle fractal, named after the Polish mathematician Wacław Sierpiński [7]. The Sierpiński triangle is one of the most simple and elegant fractals and has been a popular decorative pattern for centuries. This pattern is created by recursively each splitting each solid equilateral triangle into four smaller equilateral triangles and removing the middle one. Theoretically, these steps are repeated forever, but in practice when creating this fractal using a computer, some maximum depth must be specified as computers only have finite memory.

Figure 1.4.i Sierpiński triangle (left) [8] *and tetrahedron (right)* [8] *both of recursive depth 5*

As the recursive depth of the fractal increases, so does the number of objects (either triangles or tetrahedrons) in the scene. The total number of objects in the Sierpiński triangle increases by a factor of 3 each iteration and the tetrahedron by a factor of 4. This is the limiting factor when rendering the Sierpiński tetrahedron, as computer memory is finite and can only store limited number of objects.

2.1.2 Menger Sponge

The Menger sponge, also known as the Menger cube or Sierpiński cube, is another 3D representation of one of Sierpiński's 2D fractals. This 2D fractal is known as the Sierpiński carpet, which follows very similar recursive rules to the Sierpiński triangle but uses squares instead of triangles.

Figure 1.4.ii Sierpiński carpet (left) [9] *and Menger sponge (right)* [10] *both of recursive depth 4*

The number of objects required to create these fractals at various recursive depths increases similarly to the Sierpiński triangle and tetrahedron, but at a larger rate. The Sierpiński carpet increases by a factor of 8 each iteration while the Menger sponge by a factor of 20. A Menger sponge at recursive depth n is made up of 20^n smaller cubes. As with the Sierpiński tetrahedron, the limiting factor when trying to create this fractal is the exponential growth of the number of objects in the scene as the recursive depth increases.

In addition to the Sierpiński tetrahedron and Menger sponge, Sierpiński variations of other convex polyhedrons exist. A convex polyhedron is the 3D equivalent of a 2D regular polygon. While there are an infinite number of regular polygons, there are only fixe possible convex polyhedrons. These are the tetrahedron, cube, octahedron, dodecahedron, and icosahedron. These polyhedrons are known as the platonic solids, and their fractal counterparts are called the platonic solid fractals. This Wikipedia page contains a concise and informative list of all these shapes [11].

2.1.3   Mandel Bulb
While the platonic solid fractals are created by making many copies of a primitive shape, another method for creating fractals is to plot the convergence of values for certain mathematical equations. This can instead generate much more "natural" looking fractal patterns. One of the first fractals of this type to be discovered was the Mandelbrot set, discovered by Adrien Douady and named after Beniot Mandelbrot, the inventor of the concept of fractal geometry. The Mandelbrot fractal is defined as the set of complex numbers c for which the iteration from $z = 0$ in the equation $z_{n+1} = z_n^2 + c$ remains bounded (does not diverge to infinity) [12]. This definition is commonly written in the form $f_c(z) = z^2 + c$, where the value of c is varied. While this equation is staggeringly simple, when plotted on a

graph using the value of c (a complex number in the form $x + iy$) as the position on the axis, with x and y corresponding to the position on the x and y axis, a colour can be assigned based on how quickly that value of c tends towards infinity. This results in a beautiful shape that when zoomed in on, shows more fractal shapes within. This Wikipedia page contains a fantastic image gallery showing common shapes found within the Mandelbrot set [13], a few of which are shown below.

Figure 1.4.iii Mandelbrot set overview (left) [13]*, antenna (middle)* [13]*, and upside-down seahorse (right)* [13]

The Mandel bulb is a commonly used 3D representation of the 2D Mandelbrot fractal, created by Daniel White and Paul Nylander. For many years, it was thought that a true 3D representation of the Mandelbrot fractal did not exist, since there is no 3D representation of the 2D space of complex numbers, on which the Mandelbrot fractal is built upon [14]. While this is still the case, White and Nylander made a significant breakthrough which resulted in the creation of a 3D fractal bearing similar characteristics to the 2D Mandelbrot set.

White and Nylander considered some of the geometrical properties of the complex numbers. The multiplication of two complex numbers is a kind of rotation, and the addition is a kind of transformation. White and Nylander experimented with ways of preserving these characteristics when converting from 2D to 3D, and their solution was to change the squaring part of the formula to instead use a higher power, a practice sometimes used with the 2D Mandelbrot fractal to produce snowflake type results [15]. This change leads us to the equation $f(z) = z^n + c$, where z and c are triplex numbers, representing a point with an x, y, and z coordinate. The value of n is varied to give different results, where $n = 8$ is commonly used as this results in a good amount of fractal detail when zooming in.

White and Nylander's formula for the nth power of a point in 3D space [15], [16] is given as:

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix}^n = r^n \begin{bmatrix} \sin(n\theta)\cos(n\varphi) \\ \sin(n\theta)\sin(n\varphi) \\ \cos(n\theta) \end{bmatrix}$$

$$where\ r = \sqrt{x^2 + y^2 + z^2}, \theta = atan2(\sqrt{x^2 + y^2}, z), \varphi = atan2(y, x)$$

And the addition of two points is given as:

$$\begin{bmatrix} x_1 \\ y_1 \\ z_1 \end{bmatrix} + \begin{bmatrix} x_2 \\ y_2 \\ z_2 \end{bmatrix} = \begin{bmatrix} x_1 + x_2 \\ y_1 + y_2 \\ z_1 + z_2 \end{bmatrix}$$

Using both formulas, the equation $f(z) = z^n + c$ can now easily be solved and when rendered in 3D results in some beautiful images. Some renders from Daniel White's website are shown below.

Figure 1.4.iv Mandel bulb power of two (left) [17]*, three (middle)* [17]*, and eight (right)* [17]

The power of two version of the function (left) results in a filled 3D Mandelbrot set, with little fractal detail. Power of three (middle) contains more detail, and power of eight (right) is the sweet spot for this formula in the trade-off between fractal detail and performance. There exist several variations of Mandel bulb formula, each balancing performance with fractal detail.

The most notable performance improvement for this algorithm was discovered by David Mankin [18], and uses a distance estimation function, which for any point in 3D space, returns an estimation of the distance to the surface of the geometry. Distance estimation functions exist for many other 3D fractals as well, meaning that this is a valid generalised method for rendering 3D fractals.

In addition to the Mandel bulb, there exist many other 3D fractals. One interesting example are the Julia sets, which come from the same $f(z) = z^2 + c$ equation as the Mandelbrot set, but the value of z is varied instead of the value of c. An analysis of this fractal and will be completed in the future.

2.1.4   3D Fractals Summary

In summary, there exist two main ways of creating 3D fractals. The first approach is used for creating the platonic solid fractals and relies on taking primitive shapes and applying transformations, rotations, and scaling operations to them. The main bottleneck of rendering this type of fractal is the number of objects in the scene that must be rendered. The second approach is used for rendering more natural looking fractals and relies on plotting the convergence of equations and arbitrarily colouring them depending on how quickly the values converge. The main bottleneck of this approach is the number of iterations required for the values to converge, which can be improved using a distance estimation function which for any point in 3D space, returns an estimation of the distance to the surface of the geometry.

Both approaches discussed create completely different looking fractals, and the application will contain examples of both. However, to view these fractals, a suitable rendering approach which supports both primitive object transformation, rotation, and scaling operations, and supports the rendering of a surface defined by a series of points in 3D space. This will allow both fractal rendering approaches to be supported.

2.2   FRACTAL RENDERING METHODS

2.2.1   Rasterization

In computer graphics, there are two commonly used methods of rendering an image of a 3D scene. The first is called rasterization, which renders an image of a scene by looping through all objects in that scene, determining which pixels on the screen are affected by that object, and modifying them accordingly. This approach has many benefits [19], most notably its speed when rendering an image. Additionally, when doubling the number of pixels in an image, the time taken to rasterize the image increases by less than double. Because of these benefits, rasterization is the most common rendering method, and all graphics cards (GPUs) are designed to efficiently render images using this approach.

Rasterization is very well suited for rendering 3D objects that are stored as meshes, which contain vertices, edges and faces. Unfortunately, a 3D fractal could not be converted into a mesh without losing detail, as a mesh only contains a finite number of points, and a 3D fractal must contain infinite detail.

2.2.2 Ray Tracing
Ray tracing is another method of rendering an image of a 3D scene. When rendering an image using ray tracing, for each pixel in the camera, a ray (simply a line in 3D space) is extended or traced forwards from the camera position until it intersects with the surface of an object. From there, the ray can be absorbed or reflected by the surface and more rays can be sent out recursively. Ray tracing is ideal for photorealistic rendering as it takes into consideration many of the properties of light, through simulating reflections, light refraction, and reflections of reflections [20]. Often, ray tracers do not render images in real-time as the process is computationally expensive. To make a ray tracer capable of rendering in real-time, many approximations must be made, or hybrid approaches used.

With a little modification, the ray tracing algorithm can be modified to render an image represented by a distance estimation function instead. This approach is called ray marching, and it can be used to render 3D fractals since fractal geometry can be represented by a distance estimation function. The optimisation for the Mandel bulb fractal which uses distance estimation [18] gives such a massive performance increase when compared to ray tracing, that this approach can be used to render fractals in real time.

2.3 RAY MARCHING
Ray marching is a variation of ray tracing, which only differs in the method of detecting intersections between the ray and geometry. Instead of a using a ray-surface intersection function which returns the position of intersection, ray marching uses a distance estimation (DE) function, which simply returns the distance from any given position in the scene, to the surface of the closest geometry. Instead of shooting the ray in one go, ray marching uses an iterative approach, where the current

position is moved/marched along the ray in small increments until it lands on the surface of an object. For each point on the ray that is sampled, the DE function is called and marched forward by that distance, and the process is repeated until the ray lands on the surface of an object. If the distance function returns zero at any point (or is close enough to an arbitrary epsilon value), then then the ray has collided with the surface of the geometry. The diagram below shows a ray being marched from position p0 in the direction to the right. The distance estimation for each point is marked using the circle centred on that point.

Figure 2.3.i Ray marching diagram

Technically, the DE does not have to return the exact distance to an object, as for some objects this may not be computable, but it must never be larger than the actual value. However, if the value is too small, then the ray marching algorithm becomes inefficient, so a fine balance must be found between accuracy and efficiency.

2.3.1   Benefits of Ray Marching

Ray marching may sound more computationally complex than ray tracing since it must complete multiple iterations of an algorithm do what ray tracing does in a single ray-surface intersection function, however, it does provide several benefits. As mentioned above, ray marching does not require a surface intersection function like ray tracing does, which means it can be used to render geometry for which these functions do not exist.

While many effects such as reflections, hard shadows and depth of field can be implemented almost identically to how they are in ray tracing, there are several optical effects that the ray marching algorithm can compute very cheaply.

Ambient occlusion is a technique used to approximate how exposed each point in a scene is to ambient lighting [20]. This means that the more complex the surface of the geometry is (with creases, holes etc), the less ways ambient light can get into it those places and so the darker they should be. With ray marching, the surface complexity of geometry is usually proportional to the number of steps taken by the algorithm [21]. This property can be used to implement ambient occlusion and comes with no extra computational cost at all.

Soft shadows can also be implemented very cheaply, by keeping track of the minimum angle from the distance estimator to the point of intersection, when marching from the point of intersection towards the light source [22]. This second round of marching must be done anyway if any type of lighting is to be taken into consideration, so minimum check required for soft shadows is practically free.

A glow can also be applied to geometry very cheaply, by keeping track of the minimum distance to the geometry for each ray. Then, if the ray never actually collided with the geometry, a glow can be applied using the minimum distance the ray was from the object, a strength value, and colour specified [21].

2.3.2   Signed Distance Functions

A signed distance function (SDF) for a geometry, is a function which given any position in 3D space, will return the distance to the surface of that geometry. If the distance contains a positive sign if the position is outside of the object, and a negative sign if the position is inside of the object. If a distance function returns zero for any position, then the position must be exactly on the surface of an object. Every single geometry in a scene must have its own SDF. The scenes distance estimation (DE) function will loop through all the SDF values for the geometry in the scene and will return the closest to the point specified.

The sign returned by the SDF is useful as it allows the ray marcher to determine if a camera ray is inside of a geometry or not, and from there it can use that information to render the objects differently. We may want to render geometry either solid or hollow, or potentially add transparency.

Signed distance functions are already known for most primitive 3D shapes, such as spheres, boxes, and planes. A full list of these functions can be found on Inigo Quilez's web page [23]. A full example of the SDF for a sphere is included in appendix section 9.2. Since distance estimation can be used to represent primitive shapes, the only thing still required to render the platonic solid fractals is the ability to transform, rotate and scale primitives.

2.3.3 Transforming SDFs

Translating and rotating objects is trivial when using distance estimation. All that is required, is to transform the point being sampled with the with the inverse of the position and rotation used to place an object in the scene [23]. A function to transform a point by a translation and rotation is given below.

$$transform(p, t) = invert(t) * p$$

$$where\ p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}, t\ is\ a\ 3 \times 4\ transformation\ matrix\ storing\ only\ translation\ and\ rotation$$

In addition, uniform scaling of an SDF can be completed by first decreasing the scale back to one, sampling the point, and then increasing the scale back up again. Functions for decreasing and increasing the scale of a point in 3D space are given below.

$$scaleDown(p, s) = \frac{p}{s}, \qquad scaleUp(p, s) = p \times s, \qquad where\ p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}, s \in \mathbb{R}\ is\ the\ scale$$

2.3.4 Combining SDFs
Another feature that would be nice to include in the renderer is the ability to combine SDFs, and to have multiple objects in the same scene. SDFs can be combined using the union, subtraction, and intersection operations [24], as given below.

$$union(a, b) = \min(a, b),\ subtraction(a, b) = max(-a, b),\ intersection(a, b) = max(a, b)$$
$$where\ a, b \in \mathbb{R}\ are\ the\ values\ returned\ from\ object\ a\ and\ b's\ SDF$$

There also exist variations of formulas above which can be used to apply a smoothing value to the operation. These formulas have been included in the appendix section 9.3. The images below were rendered using an early prototype of the application.

Figure 2.3.ii Ray marched sphere and box scene experiment union (left), subtraction (middle) and smooth union (right)

There are several additional alterations that can be applied to primitives once we have their SDF [23]. A primitive can be elongated along any axis, its edges can be rounded, it can be extruded, and it can be "onioned" – a process of adding concentric layers to a shape. All these operations are relatively cheap. Signed distance functions can also be repeated, twisted, bent, and surfaces displaced using an equation such as a noise function or sin wave, though these alterations are more expensive. All these techniques mentioned will be essential when creating more complex geometry.

2.3.5   Surface Normal

The surface normal of a position on the surface of a geometry, is a normalised vector that is perpendicular to the that surface. This information is is essential for most lighting calculations and surface shading techniques, such as phong shading [20] (see appendix 9.1 for an example). Lighting and surface shading functionality has been added as stretch goals in the requirements specification section 3.1. When using distance estimation, the surface normal of any point on the geometry in a scene can be determined by probing the SDF function on the x, y and z axis, using an epsilon value of arbitrary value. The formula used to calculate the surface normal of any point using distance estimation can be found in appendix 9.4.

2.3.6   Ray Marching Summary

In summary, ray marching is a variation of ray tracing which contains all the building blocks required for rendering both types of 3D fractals. It uses distance estimation and signed distance functions for calculating intersections with geometry, primitives can be modelled using SDFs, and SDFs can be translated, rotated, and scaled. In addition, there are many mathematical operations that can be used for combining SDFs to create more complex scenes and the surface normal of geometry can be calculated which means that it is possible to create advanced rendering features like surface

shaders. The next section will give a brief introduction to GPU computing and how it will be used in the project to execute the rendering code in parallel on the GPU.

2.4 GPU COMPUTING

GPU computing is when a GPU (Graphics Processing Unit) is used in combination with a CPU (Central Processing Unit) to execute some code [25]. The correct use of GPU computing improves the overall performance of the program by offloading some computation from the CPU to GPU. A CPU is designed for executing a sequence of operations, called a thread. A CPU can execute a few tens of threads in parallel and is designed to execute them as fast as possible [26]. On the other hand, a GPU is designed to execute a few thousand of these threads in parallel but does it significantly slower than the CPU would, achieving a higher overall throughput. This difference in architecture between CPUs and GPUs means that CPUs are better suited for computations requiring data caching and flow control, while GPUs are better suited for highly parallel computations.

Implementing a graphics renderer using GPU computing is the perfect choice, as this is a massively parallel task requiring a computation to be executed for every single pixel on the screen. Making use of this architecture is the only feasible way to implement a real-time renderer. When implementing a real-time fractal renderer, the extra features that libraries like OpenGL [27] provide, such as compute shaders, aren't needed. Instead, it makes more sense to make use of general-purpose GPU computing library, as the overhead of this software will be significantly less, giving us better performance.

2.4.1 CUDA vs OpenCL

Currently, there are two suitable general-purpose GPU computing interfaces, CUDA (Compute Unified Device Architecture) and OpenCL (Open Computing Language). Both interfaces make use of a kernel language, a specialised programming language in which the code to be executed in parallel is written. The table below breaks down the main advantages and disadvantages [26], [28], [29] of CUDA and OpenCL.

Table 2.4.i CUDA and OpenCL comparison

Comparison	CUDA	OpenCL
Portability	Only runs on NVIDIA hardware	Runs on pretty much all hardware - NVIDIA, AMD, Intel, Apple, Radeon, etc
Open Source	Propriatery framework of NVIDIA	Open source standard
Technicalities		Compiles kernel code at runtime, which could take a lot of time, though this also means that the compiled code is more optimised for the device currently running it
Interface Languages	C, C++, Fortran, Java Python Wrappers, DirectCompute, Directives (e.g. OpenACC)	C, C++, Python wrapper
Kernel Language	KUDA C++	OpenCL C (C99) and OpenCL C++ (C++17)
Libraries	Extensive and powerful libraries	Good libraries, but not as extensive as CUDA
Performance	No clear advantage	No clear advantage
OS Support	Runs on Windows, Linux and MacOS	Runs on Windows, Linux and MacOS

The main difference between CUDA and OpenCL, is CUDA can only run on NVIDIA branded GPUs. While this technically does allow CUDA to make full use of the hardware, it makes the application deployment far less portable as the code will not run on other types of GPUs. OpenCL is, therefore, the superior choice when aiming to write portable GPU code.

2.5 REVIEW OF EXISTING APPLICATIONS

A review of relevant and popular existing 3D fractal renderers was completed, and the key features of each application was recorded. A summary containing points to take away from the review is included at the end.

2.5.1 Fragmentarium [30], [31]

- Features an editor within the application, which allows scenes to be created and previewed in real time

- Contains many example scenes

- Provides very detailed parameters to edit the scene

- Rendering quality of the output image is acceptable (see appendix 9.5)

22

2.5.2 Mandelbulb3D [32]

- User interface is clunky, and features are hard to find

- Only contains a couple example scenes

- The performance of the real-time preview feature is sub optimal, making it hard to navigate the scene

- The output image is very detailed, is coloured beautifully and contains many advanced optical effects (see appendix 9.6)

2.5.3 Existing Applications Summary

After reviewing existing 3D fractal rendering applications, there are several key points that should be considered when designing the application.

- The application must contain lots of example scenes, and they must be easy to find

- The real-time preview of the scene must have good performance and be easy to control

- It would be nice to be able to output an image of the scene, with an easy-to-use interface

- It must be relatively easy to create a new scene and preview it

Additionally, Fragmentarium and Mandelbulb3D are both built using OpenGL with GLSL shaders, as this makes creating scenes simple as it uses an already existing language and format. However, using OpenGL is unnecessary as not many of the features it provides are being used.

2.6 LITERATURE REVIEW SUMMARY

In summary, there are two main methods of creating fractal patterns - recursively applying transformations to primitive shapes and plotting the convergence of equations. To be able to render both types of fractals, a variation of ray tracing called ray marching must be used which uses a distance estimation function to calculate the distance to the closest piece of geometry in the scene. In addition, to be able to render 3D fractals in real time, GPU computing techniques must be used to execute code in parallel. Several existing applications have already used similar approaches to great

success, however, these applications do contain some flaws and a note of these have been made so that this project will not make the same mistakes.

3 REQUIREMENTS ANALYSIS

3.1 REQUIREMENTS SPECIFICATION

The tables below display the key functionality to be implemented to achieve the projects aim and objectives, specified in section 1.2. Requirements have been grouped by the project objective that they fall under. Requirements were prioritised using the following strategy:

- MUST – a requirement that is of the highest priority to the project

- SHOULD – a requirement that is not essential, but it would be good if the project had it

- COULD – a requirement that is optional

- WON'T – a requirement that would be implemented if the project could run for more time

Table 3.1.i Functional requirement specification

ID	Name	Description	Priority	Testing strategy
FR-1	**Objective 3**	**Core functionality**		
FR-1.1	Scene requirements	A scene must contain: • Geometry • Lights • Camera	MUST	Unit test
FR-1.2	Example scenes	The application must contain multiple example scenes, some of which should include: • Sierpinski tetrahedron fractal • Menger sponge fractal • Julia set fractal • Mandel bulb fractal	SHOULD	Unit test
FR-1.3	Mandatory optical effects	The application must support the following optical effects: • Ambient occlusion • Hard and soft shadows • Glow	MUST	Unit test
FR-1.4	Load scene	Scenes must be defined each in their own kernel file and the user must be able to load them into the application at runtime	MUST	Unit test
FR-1.5	Settings	Basic application settings must be editable. This includes:	MUST	Unit test

ID	Name	Description	Priority	Testing strategy
		• Output resolution • Optical effects enable/disable • Controls • Mouse sensitivity		
FR-2	**Objective 4**	**Additional functionality**		
FR-2.1	Real-time	The application must be capable of rendering the view of the scene in real time	MUST	Benchmark
FR-2.2	Game Loop	The application must make use of a game loop to update the scene. Methods should be called in the following order: 1. Poll for user input 2. Update scene 3. Render scene	MUST	Unit Test
FR-2.3	Controllable camera	The user must be able to control a camera, using a keyboard and mouse to move it around the scene	MUST	Unit Test
FR-2.4	Dynamic scenes	The contents of a scene (requirement FR-1.1) must be able to move around at runtime	MUST	Unit Test
FR-2.5	Optional optical effects	The application could additionally support the following optical effects: • Reflections • Depth of field • Transparency	COULD	Unit test
FR-2.6	Optional Surface Shading	The application could support the following surface shading algorithms: • Lambert • Oren-Nayar • Phong or Blinn-Phong • Subsurface scattering	COULD	Unit Test
FR-2.7	Fixed camera paths	The application camera could additionally move using a specified camera path	COULD	Unit Test
FR-2.8	Image output	It could be possible for the application to output a high-resolution image of the current camera view	COULD	Unit Test
FR-2.9	Video output	It could be possible for the application to output a video sequence for the current scene	COULD	Unit Test
FR-3	**Objective 5**	**Evaluation**		
FR-3.1	Performance benchmark	The application must contain a performance benchmark scene, as specified in section 5.2	MUST	Benchmark

Table 3.1.ii Non-functional requirement specification

ID	Name	Description	Priority	Testing strategy
NF-1	Executable	The application must run from a compiled executable	MUST	Unit test

NF-2	Operating system	The application must run on Windows 10. Where possible, cross platform libraries should be used, though other operating systems will not be officially supported	MUST	Unit test
NF-3	Display resolutions	The application must support the following common display resolutions: 1366x768, 1920x1080, 2560x1440 and 3840x2160	MUST	Unit test
NF-4	GPU Parallel Computing	The application must run in parallel on the GPU	MUST	Unit test

The requirements specified above will be implemented according to the project plan timeline specified in section 6.5. In addition, the method of evaluating the if the requirements have been successfully implemented is specified in section 5.

4 SOFTWARE DESIGN

4.1 TECHNOLOGIES
The application will be developed using the following technologies:

Table 4.1.i Application technologies

Technology	Description	Justification
OpenCL	Programming language which allows code to be run in parallel on the GPU	• GPU parallel computing gives a massive performance boost when executing the same piece of code simultaneously for many different input values • GPU parallelism is far better suited for this task than CPU parallelism as the same piece of code must be executed for every pixel on the screen • OpenCL was chosen as it has good documentation and examples, contains C and C++ programming interfaces, and allows deployments to different platforms
C++	Common system programming language	• C++ is a low-level language with good performance • C++ was chosen over C to allow an object-oriented style of programming
Microsoft Unit Testing Framework for C++	Unit testing framework	• Unit testing is essential for ensuring code is correct and functionality has been implemented • This framework comes with the Visual Studio 2019 C++ development package which is also being used
SDL2	Cross platform C++ library for manipulating windows and reading user input	• Cross platform libraries provide an abstraction layer over platform specific libraries, which allows the program implementation to remain decoupled from the deployment platform • SDL2 was chosen as it provides both window display interaction and user input event polling, and has good documentation and examples

Development of the application and documentation will be assisted the following technologies:

Table 4.1.ii Development technologies

Technology	Description	Justification
Visual Studio 2019	Code editor	• Contains powerful development tools such as refactoring options, code snippets, documentation preview etc
GitHub	Version control software	• Version control is essential for any large coding project • GitHub is being used to store all project materials, including documents, papers, and the coding project
Microsoft Word	Word processing software	• Powerful and easy to use word processing software • Documents will be edited locally and backed up to GitHub
Mendeley	Reference manager	• This reference manager has a web interface which is very convenient when researching • It also has a Microsoft Word add in which manages all referenced sources automatically
Microsoft Teams	Video communication software	• For communication with the project supervisor and second readers

4.2 CLASS STRUCTURE

The application will be structured using an object-oriented programming style to promote code reusability and data encapsulation. The application will consist of several key classes which are listed in the table below. A provisional class diagram has been included in appendix 9.7.

Table 4.2.i Class responsibilities

Class name	Responsibilities
Application	Contains the run method, the main application loop which drives the application This class contains instances of Display, Renderer and Controller
Display	Setting pixels in the display window and controlling any GUI elements
Renderer	Calculating the colour for each pixel of the display window
Input	Reading keyboard and mouse input from the user

The Display and Input classes are basic and only provide a wrapper around corresponding SDL2 window methods for setting pixels in the display, and for reading user input. The Renderer class, however, is much more complex and requires discussion. The Renderer class contains the current pixels to be displayed on the screen this frame, which is calculated using an OpenCL kernel. The OpenCL kernel is a piece of code written in one of the OpenCL kernel languages, and is the code that is executed in parallel on the GPU. Most of the ray marching code should be written in this kernel language to give the best performance to the application. Each scene will be defined within its own kernel file and will

be loaded into the application at runtime. However, this makes it hard to reuse code between kernel files as the implementations of several methods, and the values of several constants will differ between scenes. The tables in appendix 9.8 below show which methods and constants in the kernel are reusable between scenes and which are not.

A solution to reducing code duplication between the kernel files is to use the new OpenCL C++ kernel language, released in March 2021, which supports most C++17 features. The OpenCL C++ kernel language documentation can be found here [33]. The Renderer class will contain a main kernel file which contains the implementation of all reusable methods, along with an empty distance estimation method which the other kernel scene files must override, using method overriding. Method overriding is an object-oriented programming technique which allows a subclass to change the implementation of a method that a parent class is providing. This will drastically reduce the amount of code duplication between scenes as it allows the scenes to share common code. This is important for making scene creation accessible for newer users, who would be overwhelmed by a scene file containing hundreds of lines of code.

4.3 INTERFACE DESIGN

The application will feature a simple and intuitive user interface which contains a main display window, with a ribbon toolbar at the top. The toolbar will contain several tabs, including Settings, Load Scene, Create Scene, and Help. A mock user interface has been included in appendix 9.7.

5 EVALUATION STRATEGY

A goal-based evaluation strategy will be used to determine if the projects aim and objectives have been achieved. Unit tests and a benchmark scene will be used to determine how many of the requirements specified in section 3.1 have been fully implemented. For an objective to be considered achieved, all requirements prioritised MUST and SHOULD related to that objective must have been implemented. For the project aim to be considered achieved, all objectives must be achieved.

5.1 UNIT TESTING

Unit tests will be used to evaluate both the implementation status and code correctness of some requirements. These tests will be created during development of the application, often before the requirement has been implemented. This means that each requirement to be implemented can be directly linked to a series of unit tests, which makes it explicit when the functionality has been implemented correctly.

The chosen unit test framework is the Microsoft Unit Testing Framework for C++, as this comes with the Visual Studio 2019 C++ package.

5.2 PERFORMANCE BENCHMARK

A performance benchmark of the application must be implemented in order to evaluate if the real-time aspect of aim has been achieved. This section will specify the results that the benchmark scene must output, and calculations that will be completed with the data.

Table 5.2.i Results recorded from benchmark

Description	Units
Window display resolution	-
Duration of the benchmark scene	Seconds
Total number of frames rendered	-
Minimum frame time (time taken to render a frame)	Milliseconds
Maximum frame time (time taken to render a frame)	Milliseconds

Table 5.2.ii Results calculated from benchmark

Description	Units	Calculation
Average frame time	Milliseconds	$\dfrac{duration\ of\ the\ benchmark\ scene}{total\ number\ of\ frames\ rendered} \times 1000$
Average frames per second	Frames per second	$\dfrac{1000}{average\ frame\ time}$
Minimum frames per second	Frames per second	$\dfrac{1000}{minimum\ frame\ time}$
Maximum frames per second	Frames per second	$\dfrac{1000}{maximum\ frame\ time}$

In addition, the PC specs of the computer running the benchmark must be recorded. This information can be exported to a file by running the "System information" program on windows and selecting the

File>Export option. The exported file contains a list of the computer specifications and relevant driver information. The important information that will be extracted from the file is included in the table below.

Table 5.2.iii Relevant PC specification values

Description	Units	Line number in exported file
Operating system name	-	6
Processor name	-	15
Processor base clock speed	Megahertz	15
Processor number of cores	-	15
Processor number of logical cores	-	15
Total physical memory	Gigabytes	34
Total virtual memory	Gigabytes	36
Graphics card name	-	Varies, located under [Display] heading
Graphics card dedicated memory	Gigabytes	N/A
Graphics card shared memory	Gigabytes	N/A

Unfortunately, the dedicated and shared memory of the graphics card is not included in the file and must be recorded manually from viewing the task manager.

The benchmark scene has yet to be fully defined, but it must be non-trivial to render. This means it should contain multiple geometries (both fractal and primitive) and multiple lights while also making use of advanced rendering features like ambient occlusion, soft shadows, and reflections. It is important that the scene is consistent as possible between separate runs, therefore, the camera should be either stationary or move through the scene on a fixed path to view the geometries. The benchmark scene should run for a fixed duration so that it takes the same amount of time to run on all machines. The benchmark should also be run multiple times and averages taken of the results. This is to reduce negative impact of any background tasks that may be running on the machine.

6 PROJECT PLAN

6.1 PROJECT MANAGEMENT

The project will be developed using an Agile [34] approach, which uses small sprints of work to complete specific and defined tasks. This approach allows teams to respond to change quickly as requirements and plans are updated regularly. While this an individual project and the Agile approach is normally used in teams, the continuous improvement gained through this approach will be incredibly valuable for the project.

6.2 DESIGN METHODOLOGY

The application will be developed following the main object-oriented principles [35]. These include encapsulation, abstraction, inheritance, and polymorphism. An object-oriented style of programming has been chosen to promote code reusability within the application in the hopes that this will allow the kernel code to be as simple and easy to use as possible.

6.3 PROFESSIONAL, LEGAL, ETHICAL & SOCIAL ISSUES

The main principles of the open-source definition [36] state that a product must have publicly available source code, must allow modifications and derived works of the product, and allow free redistribution of the product. This project is licensed under the GNU General Public License v3.0 [37] which complies with the open-source definition and grants permission for modification, distribution, and commercial use of this product. However, all changes made to the licensed material must be documented, the modified source code must be made public, and the modified work must be distributed under the same license as this product. The license can be viewed on the project GitHub repository [6], within the LICENSE file.

The British Computer Society (BCS) codes of conduct [38] specifies the professional standards expected to be displayed by a member of the BCS, which includes working for the interest of the public, displaying professional integrity, accepting relevant authority, and showing a commitment to the profession. These standards will be respected and complied with for the duration of the project.

This project does not require any ethical approval as no studies requiring participants will be completed. Instead, the application's performance will be benchmarked over several available PCs, and results will be calculated to determine the how successful the project was.

6.4 RISK ANALYSIS

A risk analysis for the project has been completed to identify any potential risks that may affect the successfulness of the project, and a mitigation plan has been drawn up for each risk. The mitigation plan has been designed to reduce the chances of each risk happening and to reduce the negative impact of the risk if it were to happen. Each risk was then assigned a rating using the table below.

Table 6.4.i Risk rating matrix

		Severity		
		LOW	MEDIUM	HIGH
Probability	**LOW**	LOW	LOW	MEDIUM
	MEDIUM	LOW	MEDIUM	HIGH
	HIGH	MEDIUM	HIGH	HIGH

Table 6.4.ii Risk analysis matrix

ID	Description	Probability	Severity	Mitigation plan	Rating
R-1	Loss of work	LOW	HIGH	All work will be backed up regularly using online version control	MEDIUM
R-2	Change in requirements	LOW	MEDIUM	A thorough requirements specification has been prepared to reduce the probability of this happening In addition, an Agile development approach will be used, which by design minimises the negative effects of changes in requirements	LOW
R-3	Change of deadlines	LOW	HIGH	Communication from the course leaders will be regularly monitored This risk is very unlikely to happen as assurances have been made that the course deadlines will not change	MEDIUM
R-4	Delays due to learning new software	HIGH	LOW	Time was assigned during the planning stage to experiment with new software like OpenCL and GLSL shaders Additionally, some free time has been allocated at the end of the project timetable to allow for delays	MEDIUM

R-5	Delays due to illness	LOW	MEDIUM	Free time has been allocated at the end of the timetable to allow for delays	LOW
R-6	Delays due to bugs	MEDIUM	MEDIUM	Free time has been allocated at the end of the timetable to allow for delays	MEDIUM
R-7	Renderer can't be made real-time	LOW	MEDIUM	Thorough research has been completed with the conclusion that this project is feasible There is a safe core to the project and room for scaling back or adding extra functionality to the application if this were to occur In addition, performance analysis of a non-real-time renderer could be performed instead	LOW

The project risks will be regularly monitored, and mitigation of higher priority risks will be prioritised. Each sprint of work will conclude with an analysis of the current state of the project, and from there any risks likely to occur will be identified and corresponding mitigation plans will be consulted.

6.5 PROJECT TIMELINE

The project timeline for the second deliverable has been split into a total of eight sprints of work, each sprint being two weeks in duration. Sprint zero will occur in December during the Christmas holidays, and its purpose is to set up the project working environment and to refactor the existing experimentation code, created when researching and investigating existing solutions. Once this has been completed, Objective 3, the safe core of the project, will almost have been achieved. From there, sprints one and two will be used to achieve Objective 3, to add additional functionality to the application. Then, during sprint three and four, the project evaluation will be completed, and any remaining documentation will be created. Sprints five and six will be used to write up the project evaluation and complete the report, and sprint seven has been allocated as free time, in full expectation that there will be delays when completing tasks and the timeline will have to be pushed back.

A Gantt chart for the project was created using an online tool [39] and has been included below. Tasks have been coloured using the following strategy: documentation (blue), implementation (red), and unallocated (green).


Figure 6.5.i Project timeline Gantt chart deliverable two (left) and deliverable one (right)

7 CONCLUSION

The aim of this project is to develop an application which can render 3D fractal geometry in real-time. Several objectives have been defined to help guide this process. A thorough review of literature relevant to the project has been completed, which identified the two main methods of creating 3D fractals - recursively applying transformations to primitive shapes and plotting the convergence of equations. To be able to render both types of fractals, a variation of ray tracing called ray marching must be used which uses a distance estimation function to calculate the distance to the closest piece of geometry in the scene. In addition, to be able to render 3D fractals in real time, GPU computing techniques must be used to execute code in parallel on the graphics card.

Relevant existing solutions have also been analysed and their key advantages and disadvantages have been considered when creating the application design. The project has specific and measurable requirements that will be implemented over the course of many two-week sprints, as specified in the project timeline. A thorough risk analysis and investigation into professional, legal, ethical, and social issues has been completed and a strategy has been defined for evaluating how successfully the project aim will be achieved.

8 REFERENCES

[1] University of Waterloo, "Top 5 applications of fractals." https://uwaterloo.ca/math/news/top-5-applications-fractals (accessed Nov. 10, 2021).

[2] Jack Challoner, "How Mandelbrot's fractals changed the world - BBC News," 2010. https://www.bbc.co.uk/news/magazine-11564766 (accessed Nov. 10, 2021).

[3] T. Kluge, "Fractals in nature and applications," 2000. https://kluge.in-chemnitz.de/documents/fractal/node2.html (accessed Nov. 10, 2021).

[4] The Fractal Foundation, "Chapter 12 - FRACTAL APPLICATION." http://fractalfoundation.org/OFC/OFC-12-2.html (accessed Nov. 10, 2021).

[5] Wikipedia, "Fractal-generating software." https://en.wikipedia.org/wiki/Fractal-generating_software (accessed Nov. 14, 2021).

[6] S. Baarda, "3D Fractal Geometry Renderer," 2021. https://github.com/SolomonBaarda/fractal-geometry-renderer (accessed Sep. 24, 2021).

[7] Wikipedia, "Wacław Sierpiński." https://en.wikipedia.org/wiki/Wac%C5%82aw_Sierpi%C5%84ski (accessed Nov. 11, 2021).

[8] H. Segerman, "Fractals and how to make a Sierpinski Tetrahedron", Accessed: Nov. 11, 2021. [Online]. Available: http://www.segerman.org

[9] "Sierpinski Carpet." https://larryriddle.agnesscott.org/ifs/carpet/carpet.htm (accessed Nov. 11, 2021).

[10] J. Baez, "Menger Sponge | Visual Insight," 2014. https://blogs.ams.org/visualinsight/2014/03/01/menger-sponge/ (accessed Nov. 11, 2021).

[11] Wikipedia, "n-flake." https://en.wikipedia.org/wiki/N-flake (accessed Nov. 16, 2021).

[12] A. Douady, "Julia Sets and the Mandelbrot Set," *The Beauty of Fractals*, pp. 161–174, 1986, doi: 10.1007/978-3-642-61717-1_13.

[13] Wikipedia, "Mandelbrot set." https://en.wikipedia.org/wiki/Mandelbrot_set (accessed Nov. 13, 2021).

[14] V. da Silva, T. Novello, H. Lopes, and L. Velho, "Real-time Rendering of Complex Fractals," in *Ray Tracing Gems II*, 2021, pp. 529–544. doi: https://doi.org/10.1007/978-1-4842-7185-8_33.

[15] D. White, "The Unravelling of the Real 3D Mandelbrot Fractal," 2009. https://www.skytopia.com/project/fractal/mandelbulb.html (accessed Nov. 12, 2021).

[16] R. Englund, S. Seipel, and A. Hast, "Rendering Methods for 3D Fractals, Bachelor Thesis," 2010.

[17] D. White, "The Unravelling of the Real 3D Mandelbrot Fractal page 2," 2009. https://www.skytopia.com/project/fractal/2mandelbulb.html (accessed Nov. 20, 2021).

[18] D. Mankin, "True 3D mandelbrot type fractal, distance estimation optimisation," 2009. http://www.fractalforums.com/3d-fractal-generation/true-3d-mandlebrot-type-fractal/msg8073/#msg8073 (accessed Nov. 20, 2021).

[19] M. Mcguire, E. Enderton, P. Shirley, and D. Luebke, "Real-time Stochastic Rasterization on Conventional GPU Architectures," *Proceedings of the conference on high performance graphics*, pp. 173–182, 2010, Accessed: Nov. 20, 2021. [Online]. Available: http://research.nvidia.com

[20] J. Peddie, *Ray Tracing: A Tool for All*. 2019. doi: 10.1007/978-3-030-17490-3.

[21] Mikael Hvidtfeldt Christensen, "Distance Estimated 3D Fractals," 2011. http://blog.hvidtfeldts.net/index.php/2011/08/distance-estimated-3d-fractals-ii-lighting-and-coloring/ (accessed Nov. 04, 2021).

[22] I. Quilez, "Soft Shadows in Raymarched SDFs," 2010. https://iquilezles.org/www/articles/rmshadows/rmshadows.htm (accessed Nov. 04, 2021).

[23] I. Quilez, "Distance Functions," 2013. https://www.iquilezles.org/www/articles/distfunctions/distfunctions.htm (accessed Oct. 28, 2021).

[24] I. Quilez, "Distance to Fractals," 2004. https://www.iquilezles.org/www/articles/distancefractals/distancefractals.htm (accessed Nov. 04, 2021).

[25] Boston, "What Is GPU Computing?" https://www.boston.co.uk/info/nvidia-kepler/what-is-gpu-computing.aspx (accessed Nov. 16, 2021).

[26] NVIDIA, "Programming Guide :: CUDA Toolkit Documentation." https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html (accessed Nov. 17, 2021).

[27] The Khronos Group Inc, "OpenGL Overview." https://www.khronos.org/opengl/ (accessed Nov. 17, 2021).

[28] D. Exterman, "CUDA vs OpenCL: Which to Use for GPU Programming," 2021. https://www.incredibuild.com/blog/cuda-vs-opencl-which-to-use-for-gpu-programming (accessed Nov. 16, 2021).

[29] The Khronos Group Inc, "OpenCL Overview." https://www.khronos.org/opencl/ (accessed Nov. 17, 2021).

[30] "Fragmentarium (original version, now unmaintained)." https://github.com/Syntopia/Fragmentarium (accessed Nov. 17, 2021).

[31] "Fragmentarium." https://github.com/3Dickulus/FragM (accessed Nov. 17, 2021).

[32] "Mandelbulb3D." https://github.com/thargor6/mb3d#Coloring (accessed Nov. 17, 2021).

[33] Khronos®, "The C++ for OpenCL 1.0 Programming Language Documentation," 2021. https://www.khronos.org/opencl/assets/CXX_for_OpenCL.html#_the_c_for_opencl_programming_language (accessed Nov. 04, 2021).

[34] Atlassian, "What is Agile?" https://www.atlassian.com/agile (accessed Nov. 13, 2021).

[35] M. Chandel, "What are four basic principles of Object Oriented Programming?" https://medium.com/@cancerian0684/what-are-four-basic-principles-of-object-oriented-programming-645af8b43727 (accessed Nov. 13, 2021).

[36] Open Source Initiative, "The Open Source Definition," 2007. https://opensource.org/osd (accessed Nov. 18, 2021).

[37] Open Source Initiative, "GNU General Public License version 3," 2007. https://opensource.org/licenses/GPL-3.0 (accessed Nov. 18, 2021).

[38] British Computer Society, "Code of Conduct for BCS Members," 2021. Accessed: Nov. 18, 2021. [Online]. Available: https://www.bcs.org/media/2211/bcs-code-of-conduct.pdf

[39] TeamGantt, "Online Gantt Chart Maker." https://www.teamgantt.com/ (accessed Nov. 18, 2021).

[40] I. Quilez, "Mandelbulb," 2009. https://www.iquilezles.org/www/articles/mandelbulb/mandelbulb.htm (accessed Nov. 04, 2021).

[41] MockFlow, "Online Wireframe Tools, Prototyping Tools, Online Whiteboard, Design tool, UI Mockups, UX Suite, Remote design collaboration, UX Planning." https://www.mockflow.com/ (accessed Nov. 18, 2021).

9 APPENDICES

9.1 EXPERIMENTATION RENDERS

The images below were rendered using a very early version of the application, during the research and experimentation stage.

Figure 9.1.i Render of the Mandel bulb fractal (left) and cross section (right) using equation from [40]

The value of a surface normal can be converted to a colour by mapping the x, y, and z values to r, g, and b, which can be useful when debugging. The images below show the surface normals of the sphere and box experiment scene, and an example phong shading, a surface shading technique.

Figure 9.1.ii Surface normal visualised (left) and phong shading experiment (right)

9.2 SPHERE SDF EXAMPLE

Included is the signed distance function (SDF) for a sphere with radius R, positioned on the origin.

$$sphereSDF(p) = |p| - R$$

$where\ p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}, |p|\ is\ the\ magnitude\ of\ the\ vector\ p, R\ is\ the\ circle\ radius\ in\ world\ units$

Figure 9.2.i Sphere SDF diagram

9.3 SMOOTH SDF COMBINATIONS

The following equations can be used to combine two SDF values using a smoothing value.

$$smoothUnion(a, b, s) = \min(a, b) - h^2 \times \frac{0.25}{k}$$

$where\ a, b\ \in \mathbb{R}, s\ \in \mathbb{R}\ is\ the\ smoothing\ value, h = \max(s - abs(a - b), 0)$

$$smoothSubtraction(a, b, s) = \max(-a, b) + h^2 \times \frac{0.25}{k}$$

$where\ a, b\ \in \mathbb{R}, s\ \in \mathbb{R}\ is\ the\ smoothing\ value, h = \max(s - abs(-a - b), 0)$

$$smoothIntersection(a, b, s) = \max(a, b) + h^2 \times \frac{0.25}{k}$$

$where\ a, b\ \in \mathbb{R}, s\ \in \mathbb{R}\ is\ the\ smoothing\ value, h = \max(s - abs(a - b), 0)$

9.4 SDF Surface Normal Calculations

$$normal = normalise(\begin{bmatrix} DE(p + \begin{bmatrix} e \\ 0 \\ 0 \end{bmatrix}) - DE(p - \begin{bmatrix} e \\ 0 \\ 0 \end{bmatrix}) \\ DE(p + \begin{bmatrix} 0 \\ e \\ 0 \end{bmatrix}) - DE(p - \begin{bmatrix} 0 \\ e \\ 0 \end{bmatrix}) \\ DE(p + \begin{bmatrix} 0 \\ 0 \\ e \end{bmatrix}) - DE(p - \begin{bmatrix} 0 \\ 0 \\ e \end{bmatrix}) \end{bmatrix})$$

$$where\ p = \begin{bmatrix} x \\ y \\ z \end{bmatrix}, e\ is\ the\ epsilon\ value, DE\ is\ the\ scene\ distance\ estimation\ function$$

9.5 Fragmentarium Renders

The following screenshots were taken when experimenting with Fragmentarium [30], [31].

Figure 9.5.i A recursive scene (left) and Mandel bulb (right)

Figure 9.5.ii Tree fractal

9.6 MANDELBULB3D RENDERS

The screenshot below was rendered using MandelBulb3D [32].

Figure 9.6.i Mandel bulb fractal with natural looking colouring
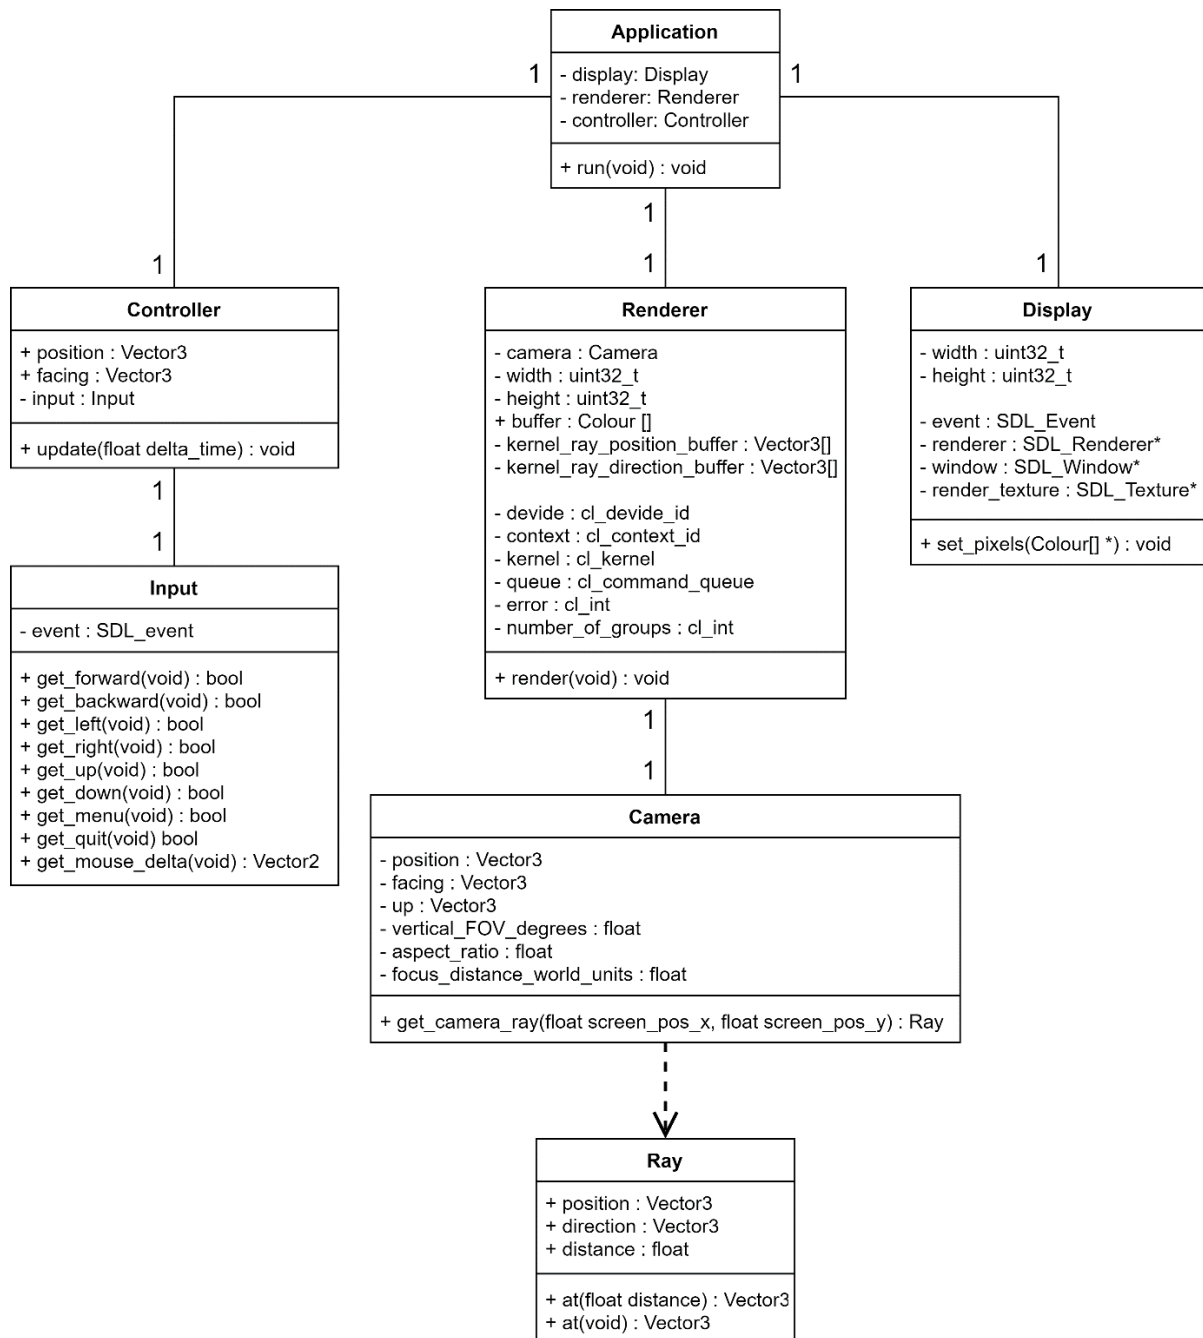
9.7   APPLICATION CLASS DIAGRAM

Application

1 | - display: Display
- renderer: Renderer
- controller: Controller | 1

+ run(void) : void

1

1   1   1

Controller

+ position : Vector3
+ facing : Vector3
- input : Input

+ update(float delta_time) : void

1

1

Input

- event : SDL_event

+ get_forward(void) : bool
+ get_backward(void) : bool
+ get_left(void) : bool
+ get_right(void) : bool
+ get_up(void) : bool
+ get_down(void) : bool
+ get_menu(void) : bool
+ get_quit(void) bool
+ get_mouse_delta(void) : Vector2

Renderer

- camera : Camera
- width : uint32_t
- height : uint32_t
+ buffer : Colour []
- kernel_ray_position_buffer : Vector3[]
- kernel_ray_direction_buffer : Vector3[]

- devide : cl_devide_id
- context : cl_context_id
- kernel : cl_kernel
- queue : cl_command_queue
- error : cl_int
- number_of_groups : cl_int

+ render(void) : void

1

1

Camera

- position : Vector3
- facing : Vector3
- up : Vector3
- vertical_FOV_degrees : float
- aspect_ratio : float
- focus_distance_world_units : float

+ get_camera_ray(float screen_pos_x, float screen_pos_y) : Ray

Display

- width : uint32_t
- height : uint32_t

- event : SDL_Event
- renderer : SDL_Renderer*
- window : SDL_Window*
- render_texture : SDL_Texture*

+ set_pixels(Colour[] *) : void

Ray

+ position : Vector3
+ direction : Vector3
+ distance : float

+ at(float distance) : Vector3
+ at(void) : Vector3

Figure 9.7.i Application class diagram

9.8 APPLICATION SCENE KERNEL REUSABILITY

Table 9.8.i Kernel method reusability matrix

Method name	Purpose	Reusable across scenes?
render()	Calculates the colour for all pixels in the display and puts the values into a buffer	YES
calculatePixelColour(Ray)	Calculates the colour for the camera pixel with the specified ray direction	YES
DE(Vector3)	Calculates the distance to the nearest geometry surface in the current scene	NO
calculateNormal(Vector3)	Calculates the surface normal vector of the geometry for the position specified	YES

Table 9.8.ii Kernel constant reusability matrix

Constant name	Purpose	Reusable across scenes?
MAXIMUM_MARCH_STEPS	Maximum number of iterations the ray marching algorithm can make	NO
MAXIMUM_MARCH_DISTANCE	Maximum distance the ray can be marched in the scene	NO
SURFACE_INTERSECTION_EPSILON	A very small value used to determine when the DE has converged to zero	YES
SURFACE_NORMAL_EPSILON	Arbitrary distance to probe the DE function when calculating the surface normal	YES

9.9 MOCK APPLICATION USER INTERFACE
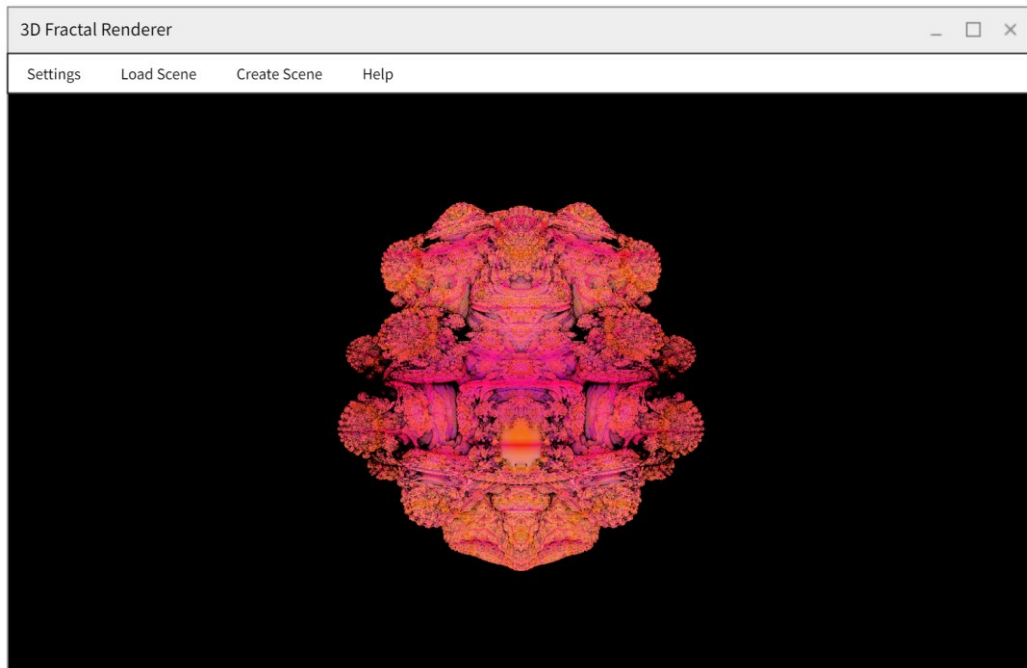
A mock user interface for the application was created using an online tool [41].

Figure 9.9.i Mock Application User Interface